

# Movie Recommendation System

Silas Mwandikwa

2023-08-30

## R Markdown

We will be building a movie recommendation system for movie sites or applications such as Netflix

We load the libraries first

```
library(recommenderlab)
```

```
## Loading required package: Matrix
```

```
## Loading required package: arules
```

```
##
```

```
## Attaching package: 'arules'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      abbreviate, write
```

```
## Loading required package: proxy
```

```
##
```

```
## Attaching package: 'proxy'
```

```
## The following object is masked from 'package:Matrix':
```

```
##
```

```
##      as.matrix
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      as.dist, dist
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
##      as.matrix
```

```
## Registered S3 methods overwritten by 'registry':
```

```
##      method          from
```

```
##      print.registry_field proxy
```

```
##      print.registry_entry proxy
```

```
library(data.table)
library(reshape2)
```

```
##
## Attaching package: 'reshape2'

## The following objects are masked from 'package:data.table':
##
##      dcast, melt
```

```
library(ggplot2)
```

## Loading the data

```
movie_data <- read.csv("IMDB-Dataset/movies.csv",stringsAsFactors = FALSE) #Letting strings remain strings
rating_data <- read.csv("IMDB-Dataset/ratings.csv")
str(movie_data)
```

```
## 'data.frame': 10329 obs. of 3 variables:
## $ movieId: int 1 2 3 4 5 6 7 8 9 10 ...
## $ title : chr "Toy Story (1995)" "Jumanji (1995)" "Grumpier Old Men (1995)" "Waiting to Exhale (1995)" ...
## $ genres : chr "Adventure|Animation|Children|Comedy|Fantasy" "Adventure|Children|Fantasy" "Comedy|Drama|Fantasy|Romance" ...
```

```
summary(movie_data)
```

```
##      movieId      title      genres
## Min.   :      1  Length:10329  Length:10329
## 1st Qu.: 3240   Class :character Class :character
## Median : 7088   Mode  :character Mode  :character
## Mean   : 31924
## 3rd Qu.: 59900
## Max.   :149532
```

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr    1.5.0
## v lubridate  1.9.2      v tibble     3.2.1
## v purrr      1.0.1      v tidyr      1.3.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::between()      masks data.table::between()
## x tidyr::expand()       masks Matrix::expand()
## x dplyr::filter()       masks stats::filter()
## x dplyr::first()        masks data.table::first()
## x lubridate::hour()     masks data.table::hour()
## x lubridate::isoweek()  masks data.table::isoweek()
## x dplyr::lag()          masks stats::lag()
```

```
## x dplyr::last()      masks data.table::last()
## x lubridate::mday()  masks data.table::mday()
## x lubridate::minute() masks data.table::minute()
## x lubridate::month() masks data.table::month()
## x tidyr::pack()      masks Matrix::pack()
## x lubridate::quarter() masks data.table::quarter()
## x dplyr::recode()    masks arules::recode()
## x lubridate::second() masks data.table::second()
## x purrr::transpose() masks data.table::transpose()
## x tidyr::unpack()    masks Matrix::unpack()
## x lubridate::wday()   masks data.table::wday()
## x lubridate::week()   masks data.table::week()
## x lubridate::yday()   masks data.table::yday()
## x lubridate::year()   masks data.table::year()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
glimpse(movie_data)
```

```
## Rows: 10,329
## Columns: 3
## $ movieId <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, ~
## $ title   <chr> "Toy Story (1995)", "Jumanji (1995)", "Grumpier Old Men (1995)~
## $ genres  <chr> "Adventure|Animation|Children|Comedy|Fantasy", "Adventure|Chil~
```

```
head(movie_data)
```

```
##   movieId      title
## 1      1      Toy Story (1995)
## 2      2      Jumanji (1995)
## 3      3      Grumpier Old Men (1995)
## 4      4      Waiting to Exhale (1995)
## 5      5      Father of the Bride Part II (1995)
## 6      6      Heat (1995)
##           genres
## 1 Adventure|Animation|Children|Comedy|Fantasy
## 2      Adventure|Children|Fantasy
## 3      Comedy|Romance
## 4      Comedy|Drama|Romance
## 5      Comedy
## 6      Action|Crime|Thriller
```

```
tail(movie_data)
```

```
##   movieId      title      genres
## 10324 146656      Creed (2015)      Drama
## 10325 146684      Cosmic Scrat-tastrophe (2015) Animation|Children|Comedy
## 10326 146878      Le Grand Restaurant (1966)      Comedy
## 10327 148238      A Very Murray Christmas (2015)      Comedy
## 10328 148626      The Big Short (2015)      Drama
## 10329 149532      Marco Polo: One Hundred Eyes (2015)      (no genres listed)
```

```
summary(rating_data)
```

```
##      userId      movieId      rating      timestamp
## Min.   : 1.0    Min.   : 1    Min.   :0.500    Min.   :8.286e+08
## 1st Qu.:192.0  1st Qu.: 1073   1st Qu.:3.000    1st Qu.:9.711e+08
## Median :383.0  Median : 2497   Median :3.500    Median :1.115e+09
## Mean   :364.9  Mean   : 13381   Mean   :3.517    Mean   :1.130e+09
## 3rd Qu.:557.0  3rd Qu.: 5991   3rd Qu.:4.000    3rd Qu.:1.275e+09
## Max.   :668.0  Max.   :149532   Max.   :5.000    Max.   :1.452e+09
```

```
movie_genre <- as.data.frame(movie_data$genres, stringsAsFactors = FALSE)
# We are combining movie_data and the genres to one data frame as strings
library(data.table)
movie_genre2 <- as.data.frame(tstrsplit(movie_genre[,1], '[ | ]',
                                     type.convert = TRUE),
                             stringsAsFactors=FALSE)
# Splitting the strings in the movie_genre df we just created, and ensuring strings
#aren't changed to factors
colnames(movie_genre2) <- c(1:10) # We'll have 10 columns
glimpse(movie_genre2)
```

```
## Rows: 10,329
## Columns: 10
## $ '1' <chr> "Adventure", "Adventure", "Comedy", "Comedy", "Comedy", "Action",~
## $ '2' <chr> "Animation", "Children", "Romance", "Drama", NA, "Crime", "Romanc~
## $ '3' <chr> "Children", "Fantasy", NA, "Romance", NA, "Thriller", NA, NA, NA,~
## $ '4' <chr> "Comedy", NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ '5' <chr> "Fantasy", NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ '6' <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, N~
## $ '7' <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, N~
## $ '8' <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, N~
## $ '9' <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, N~
## $ '10' <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, N~
```

```
list_genre <- c("Action", "Adventure", "Animation", "Children", "Comedy",
               "Crime", "Documentary", "Drama", "Fantasy", "Film-Noir", "Horror",
               "Musical", "Mystery", "Romance", "Sci-Fi", "Thriller", "War",
               "Western")
```

```
genre_mat1 <- matrix(0,10330,18) # create a matrix based on our number of rows and #cols, remember the
genre_mat1[1, ] <- list_genre #ensuring to make it two dimensional with a focus on #first column
colnames(genre_mat1) <- list_genre
```

```
for (index in 1:nrow(movie_genre2)) {
  for (col in 1:ncol(movie_genre2)) {
    gen_col = which(genre_mat1[1, ] == movie_genre2[index,col])
    genre_mat1[index+1,gen_col] <- 1
  }
}
```

```
# Basically looking at the data from both movie_genre2 and genre_mat1 and comparing #what matches in bo
```

```
genre_mat2 <- as.data.frame(genre_mat1[-1, ], stringsAsFactors=FALSE)
```

```
#Remove first row, which was the genre list
```

```
for (col in 1:ncol(genre_mat2)) {
  genre_mat2[,col] <- as.integer(genre_mat2[,col]) #convert from characters to integers
}
```

```
str(genre_mat2)
```

```
## 'data.frame': 10329 obs. of 18 variables:
## $ Action : int 0 0 0 0 0 1 0 0 1 1 ...
## $ Adventure : int 1 1 0 0 0 0 0 1 0 1 ...
## $ Animation : int 1 0 0 0 0 0 0 0 0 0 ...
## $ Children : int 1 1 0 0 0 0 0 1 0 0 ...
## $ Comedy : int 1 0 1 1 1 0 1 0 0 0 ...
## $ Crime : int 0 0 0 0 0 1 0 0 0 0 ...
## $ Documentary: int 0 0 0 0 0 0 0 0 0 0 ...
## $ Drama : int 0 0 0 1 0 0 0 0 0 0 ...
## $ Fantasy : int 1 1 0 0 0 0 0 0 0 0 ...
## $ Film-Noir : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Horror : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Musical : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Mystery : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Romance : int 0 0 1 1 0 0 1 0 0 0 ...
## $ Sci-Fi : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Thriller : int 0 0 0 0 0 1 0 0 0 1 ...
## $ War : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Western : int 0 0 0 0 0 0 0 0 0 0 ...
```

```
SearchMatrix <- cbind(movie_data[,1:2], genre_mat2[]) #Basically binding data in #movie_data df with a
head(SearchMatrix)
```

```
##   movieId      title Action Adventure Animation
## 1      1      Toy Story (1995)      0      1      1
## 2      2      Jumanji (1995)      0      1      0
## 3      3      Grumpier Old Men (1995)      0      0      0
## 4      4      Waiting to Exhale (1995)      0      0      0
## 5      5      Father of the Bride Part II (1995)      0      0      0
## 6      6      Heat (1995)      1      0      0
##   Children Comedy Crime Documentary Drama Fantasy Film-Noir Horror Musical
## 1      1      1      0      0      0      1      0      0      0
## 2      1      0      0      0      0      1      0      0      0
## 3      0      1      0      0      0      0      0      0      0
## 4      0      1      0      0      1      0      0      0      0
## 5      0      1      0      0      0      0      0      0      0
## 6      0      0      1      0      0      0      0      0      0
##   Mystery Romance Sci-Fi Thriller War Western
## 1      0      0      0      0      0      0
## 2      0      0      0      0      0      0
## 3      0      1      0      0      0      0
## 4      0      1      0      0      0      0
## 5      0      0      0      0      0      0
## 6      0      0      0      1      0      0
```

*#As you can see from the matrix created for instance first output of Toy Story, it #now tells us exactly*

```
ratingMatrix <- dcast(rating_data, userId~movieId, value.var = 'rating',
                      na.rm=FALSE)
ratingMatrix <- as.matrix(ratingMatrix[,-1]) #remove userids, coz they are in the
#first column, so minus 1
#Convert rating matrix into a recommenderlab sparse matrix
ratingMatrix <- as(ratingMatrix, "realRatingMatrix")
ratingMatrix
```

```
## 668 x 10325 rating matrix of class 'realRatingMatrix' with 105339 ratings.
```

```
recommendation_model <- recommenderRegistry$get_entries(dataType = "realRatingMatrix")
names(recommendation_model)
```

```
## [1] "HYBRID_realRatingMatrix"      "ALS_realRatingMatrix"
## [3] "ALS_implicit_realRatingMatrix" "IBCF_realRatingMatrix"
## [5] "LIBMF_realRatingMatrix"      "POPULAR_realRatingMatrix"
## [7] "RANDOM_realRatingMatrix"      "RERECOMMEND_realRatingMatrix"
## [9] "SVD_realRatingMatrix"        "SVDF_realRatingMatrix"
## [11] "UBCF_realRatingMatrix"
```

```
lapply(recommendation_model, "[", "description")
```

```
## $HYBRID_realRatingMatrix
## [1] "Hybrid recommender that aggregates several recommendation strategies using weighted averages."
##
## $ALS_realRatingMatrix
## [1] "Recommender for explicit ratings based on latent factors, calculated by alternating least squares."
##
## $ALS_implicit_realRatingMatrix
## [1] "Recommender for implicit data based on latent factors, calculated by alternating least squares."
##
## $IBCF_realRatingMatrix
## [1] "Recommender based on item-based collaborative filtering."
##
## $LIBMF_realRatingMatrix
## [1] "Matrix factorization with LIBMF via package recosystem (https://cran.r-project.org/web/packages/recosystem)"
##
## $POPULAR_realRatingMatrix
## [1] "Recommender based on item popularity."
##
## $RANDOM_realRatingMatrix
## [1] "Produce random recommendations (real ratings)."
##
## $RERECOMMEND_realRatingMatrix
## [1] "Re-recommends highly rated items (real ratings)."
##
## $SVD_realRatingMatrix
## [1] "Recommender based on SVD approximation with column-mean imputation."
##
```

```
## $SVDF_realRatingMatrix
## [1] "Recommender based on Funk SVD with gradient descend (https://sifter.org/~simon/journal/20061211)"
##
## $UBCF_realRatingMatrix
## [1] "Recommender based on user-based collaborative filtering."
```

```
recommendation_model$IBCF_realRatingMatrix$parameters
```


```
## $k
## [1] 30
##
## $method
## [1] "cosine"
##
## $normalize
## [1] "center"
##
## $normalize_sim_matrix
## [1] FALSE
##
## $alpha
## [1] 0.5
##
## $na_as_zero
## [1] FALSE
```

```
similarity_mat <- similarity(ratingMatrix[1:4, ],
                             method = "cosine",
                             which = "users")
#Creating a similarity matrix based on what different users watch and the films themselves
as.matrix(similarity_mat)
```

```
##           1           2           3           4
## 1          NA 0.9880430 0.9820862 0.9957199
## 2 0.9880430          NA 0.9962866 0.9687126
## 3 0.9820862 0.9962866          NA 0.9944484
## 4 0.9957199 0.9687126 0.9944484          NA
```

```
image(as.matrix(similarity_mat), main = "User's Similarities")
```

## User's Similarities



```
movie_similarity <- similarity(ratingMatrix[, 1:4], method = "cosine",
                             which = "items")
```


```
as.matrix(movie_similarity)
```

```
##           1           2           3           4
## 1          NA 0.9834866 0.9779671 0.9550638
## 2 0.9834866          NA 0.9829378 0.9706208
## 3 0.9779671 0.9829378          NA 0.9932438
## 4 0.9550638 0.9706208 0.9932438          NA
```

```
image(as.matrix(movie_similarity), main = "Movies Similarity")
```



## Movies Similarity



```
rating_values <- as.vector(ratingMatrix@data) # extracting unique ratings
unique(rating_values)
```

```
## [1] 0.0 5.0 4.0 3.0 4.5 1.5 2.0 3.5 1.0 2.5 0.5
```

```
#Creating a table of ratings to display our unique values
Table_of_Ratings <- table(rating_values) # Creating a count of movie ratings
Table_of_Ratings
```

```
## rating_values
##      0      0.5      1      1.5      2      2.5      3      3.5      4      4.5
## 6791761  1198   3258   1567   7943   5484   21729  12237  28880   8187
##      5
##   14856
```

```
library(ggplot2)
movie_views <- colCounts(ratingMatrix) # Count the views for each film
table_views <- data.frame(movie = names(movie_views),
                          views = movie_views) # Create data frame for views
table_views <- table_views[order(table_views$views,
                                decreasing = TRUE), ] # Sorting by the number of views from largest to


table_views$title <- NA
for (index in 1:10325) {
```

```
table_views[index,3] <- as.character(subset(movie_data,
                                           movie_data$movieId ==
                                           table_views[index,1])$title)
}
table_views[1:6,] #Pulp fiction is the movie with most views, no surprise there
```

```
##      movie views      title
## 296    296    325      Pulp Fiction (1994)
## 356    356    311      Forrest Gump (1994)
## 318    318    308  Shawshank Redemption, The (1994)
## 480    480    294      Jurassic Park (1993)
## 593    593    290  Silence of the Lambs, The (1991)
## 260    260    273 Star Wars: Episode IV - A New Hope (1977)
```


```
ggplot(table_views[1:6, ], aes(x = title, y = views)) +
  geom_bar(stat = "identity", fill = 'steelblue') +
  geom_text(aes(label = views), vjust=-0.3, size=3.5) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +

  ggtitle("Total Views of the Top Films")
```



```
image(ratingMatrix[1:25, 1:25], axes = FALSE, main = "Heatmap of the First 25 rows and columns")
```

## Heatmap of the First 25 rows and columns




```
movie_ratings <- ratingMatrix[rowCounts(ratingMatrix) > 50,  
                              colCounts(ratingMatrix) > 50]
```

```
movie_ratings
```

```
## 420 x 447 rating matrix of class 'realRatingMatrix' with 38341 ratings.
```

```
minimum_movies <- quantile(rowCounts(movie_ratings), 0.98)  
minimum_users <- quantile(colCounts(movie_ratings), 0.98)  
image(movie_ratings[rowCounts(movie_ratings) > minimum_movies,  
        colCounts(movie_ratings) > minimum_users],  
      main = "Heatmap of the top Users and Movies")
```

## Heatmap of the top Users and Movies




```
average_ratings <- rowMeans(movie_ratings)
qplot(average_ratings, fill=I("steelblue"), col=I("red")) +
  ggtitle("Distribution of the average rating per user")
```

```
## Warning: 'qplot()' was deprecated in ggplot2 3.4.0.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

Distribution of the average rating per user

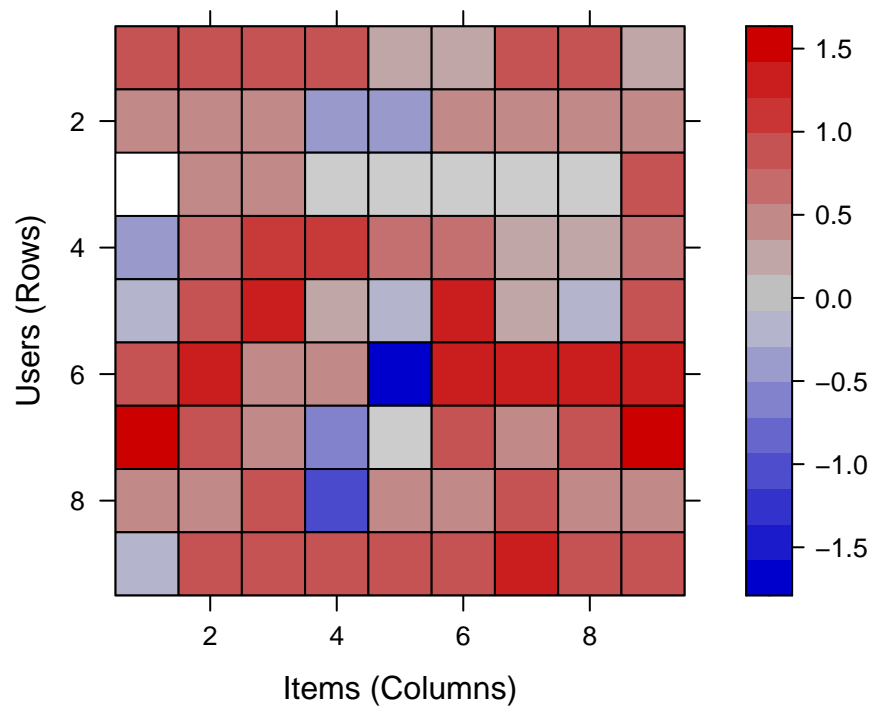


```
normalized_ratings <- normalize(movie_ratings)
sum(rowMeans(normalized_ratings) > 0.00001)
```

```
## [1] 0
```

```
image(normalized_ratings[rowCounts(normalized_ratings) > minimum_movies,
                           colCounts(normalized_ratings) > minimum_users],
      main = "Normalized Ratings for Top Users")
```

## Normalized Ratings for Top Users




Dimensions: 9 x 9

```
binary_minimum_movies <- quantile(rowCounts(movie_ratings), 0.95)
binary_minimum_users <- quantile(rowCounts(movie_ratings), 0.95)

goodRatedFilms <- binarize(movie_ratings, minRating = 3)
image(goodRatedFilms[rowCounts(movie_ratings) > binary_minimum_movies,
  rowCounts(movie_ratings) > binary_minimum_users],
  main = "Heatmap of the top users and movies")
```

## Heatmap of the top users and movies



**Dimensions: 21 x 21**

```
sampld_data <- sample(x = c(TRUE, FALSE),
                      size = nrow(movie_ratings),
                      replace = TRUE,
                      prob = c(0.8, 0.2))
```

```
trained_data <- movie_ratings[sampld_data, ]
testing_data <- movie_ratings[!sampld_data, ]
```

```
recommendation_system <- recommenderRegistry$get_entries(dataType
                                                         ="realRatingMatrix")
recommendation_system$IBCF_realRatingMatrix$parameters
```

```
## $k
## [1] 30
##
## $method
## [1] "cosine"
##
## $normalize
## [1] "center"
##
## $normalize_sim_matrix
## [1] FALSE
##
## $alpha
```

```

## [1] 0.5
##
## $na_as_zero
## [1] FALSE

recommen_model <- Recommender(data = trained_data,
                              method = "IBCF",
                              parameter = list(k = 30))
recommen_model

## Recommender of type 'IBCF' for 'realRatingMatrix'
## learned using 346 users.

class(recommen_model)

## [1] "Recommender"
## attr(,"package")
## [1] "recommenderlab"

model_info <- getModel(recommen_model)
class(model_info$sim)

## [1] "dgCMatrix"
## attr(,"package")
## [1] "Matrix"

dim(model_info$sim)


## [1] 447 447

top_items <- 20
image(model_info$sim[1:top_items, 1:top_items],
      main = "Heatmap of the first rows and columns")

```



## Heatmap of the first rows and columns



**Dimensions: 20 x 20**


```
sum_rows <- rowSums(model_info$sim > 0)
table(sum_rows)
```

```
## sum_rows
## 30
## 447
```

```
sum_cols <- colSums(model_info$sim > 0)
qplot(sum_cols, fill=I("steelblue"), col=I("red")) +
  ggtitle("Distribution of Column Count")
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

Distribution of Column Count



```
top_recommendations <- 10 #the number of items to recommend to each user
predicted_recommendations <- predict(object = recommen_model,
                                     newdata = testing_data,
                                     n = top_recommendations)
```

```
predicted_recommendations
```

```
## Recommendations as 'topNList' with n = 10 for 74 users.
```

```
user1 <- predicted_recommendations@items[[1]] #recommendation for the first user
movies_user1 <- predicted_recommendations@itemLabels[user1]
movies_user2 <- movies_user1

for (index in 1:10) {
  movies_user2[index] <- as.character(subset(movie_data,
                                             movie_data$movieId ==
                                             movies_user1[index])$title)
}
```

```
movies_user2
```

```
## [1] "Toy Story (1995)"
## [2] "Casino (1995)"
## [3] "Sense and Sensibility (1995)"
## [4] "Leaving Las Vegas (1995)"
```

```
## [5] "Seven (a.k.a. Se7en) (1995)"
## [6] "Taxi Driver (1976)"
## [7] "Like Water for Chocolate (Como agua para chocolate) (1992)"
## [8] "Léon: The Professional (a.k.a. The Professional) (Léon) (1994)"
## [9] "Blade Runner (1982)"
## [10] "Trainspotting (1996)"

recommendation_matrix <- sapply(predicted_recommendations@items,
                                function(x){ as.integer(colnames(movie_ratings)[x]))} #matrix with recommendations
recommendation_matrix[,1:4]

##           0    1     2     3
## [1,]      1    3  1674  3175
## [2,]     16   21  1704  5989
## [3,]     17  161  2355  1250
## [4,]     25  185  72998  1282
## [5,]     47  235    588  1617
## [6,]    111  349   7147  1219
## [7,]    265  357   2000   908
## [8,]    293  440    150  4011
## [9,]    541  474   4973  1214
## [10,]   778  661    555  3996
```